

Honesty Paper Pitch

Puneh Nejati-Mehr, Korbinian Friedl

The Eliciting Latent Knowledge Problem is the problem of training AI systems to honestly respond to a human interlocutor's questions.

In their 2026 paper “The Impossibility of Eliciting Latent Knowledge”, Friedl, Ward et al. formalise this problem and what makes it difficult. The notion of “honesty” employed in this paper is very limited, however. In the context of their analysis, an “honest” answer to a question of the form “What is the value of variable V ?” is the answer which states that value (in the domain of V) to which the agent, subjectively, assigns the highest probability.

This will in general be the honest response to a question asking for an agent's best guess. It will not, however, in general be the honest response asking for the value of the variable. Intuitively, a minimal requirement for the latter question to be answered honestly would be for the agent to communicate something they themselves believe. It is easy to see how “belief” and “best guess” can come apart [explain how].

We therefore want to supplement the account from Friedl and Ward (2026) by developing a more philosophically plausible account of what exactly it is that makes an assertion (in particular: a response to a question one is asked) “honest” or “dishonest”, and showing how it can be modelled in the formal framework of Causal Influence Diagrams.

The paper correspondingly has two main parts:

In the first part, we survey the philosophical literature and develop an account of the honesty of an assertion (to be distinguished, e.g., from honesty as a character trait or virtue). (For a preliminary review of the philosophical literature, see Section 1.)

The account we defend will build on the intuition that the crucial factor for an assertion to be honest is that it be performed with the intention that one's interlocutor should know the truth. This account can explain how it is possible for an agent's assertion to

- **represent their best guess for the value of a variable and still be dishonest**

The best guess may still be subjectively very unlikely, and for this (or other; think e.g. something like Buchak's “belief, credence and norms”) reasons you may not yourself believe it. If you then

communicate this content, it would not be honest, even though it is your best guess.

- **state a true description of reality which they themselves believe and still be dishonest**

Example: cheating spouse being asked by their partner “did you meet X yesterday evening?” and responding with “I had a meeting with my boss yesterday evening!”, which is true and believed by them, but carries the wrong and disbelieved implication that they didn’t also meet X.

- **state a false description of reality which they themselves disbelieve and still be honest**

Example: It is part of the common ground between agents A and B that the implication $p \rightarrow q$ holds. In this case, agent A can communicate q to B by communicating p . A can thus honestly inform B of the truth of q by communicating p , even if p is wrong and disbelieved by A.

Our account is going to center on the structure of the intention with which an agent performs the speech act of making an assertion. We are going to argue that...

Then, in the second part, we are going to use (ward2024reasons?)’s formalisation of intention in the framework of Causal Influence Diagrams to translate this account of honesty into the formal framework Friedl and Ward proposed for studying ELK.

1 Preliminary (mini) review of the philosophical literature

There is widespread agreement that honest language use is necessary for ‘good’ language use. Yet, while the philosophical and linguistic literature have given much attention to defining ‘insincere’ language use—that is, deceptive language, like lying and misleading—little attention has been given to explicitly defining a notion of honest language use (that is, outside the virtue-theoretic literature that defines honesty as a virtue of an agent and not as a property of an utterance). Furthermore, the explicit and implicit assumptions about what ‘honesty’ means cannot all be subsumed under one common understanding. More specifically, the distinctions between honesty and neighbouring concepts like truthfulness, accuracy, and lack of deception need more attention.

One disparity concerns the relation between truth-aiming (stating true beliefs) and honest language use. It is a commonly held view that ‘good’ or ‘cooperative’ communication necessitates some sense of honesty. For instance, in explaining the mechanisms of successful communication, the Gricean Cooperative Principle and more specifically the Maxim of Quality (“Try to make your contribution one that is true” (Grice 1975, 46)) seems to implicitly rely on a notion of honesty that constitutes truthful language use. But this does not suffice for an adequate notion of honesty as it enmeshes two distinct norms: an evidential norm (forming true beliefs that one would assert), and an honesty-norm (only saying what one believes). It is easy to see how these two can come apart. For instance, an agent may be always only saying what she believes to be true, but be very unreliable in forming true beliefs. It is questionable to

claim that her reliable honesty sufficiently constitutes truthfulness.

This distinction between truth-aiming and honesty is appreciated by Bernard Williams' characterisation of truthful language use as crucial to explaining successful communication. Williams (2002) distinguishes two constituent virtues of truthful language use: Accuracy and sincerity. Accuracy concerns the disposition to reliably (or rationally) form beliefs that are true (2002, 125–27), and sincerity (what we've been calling 'honesty') concerns the disposition to only assert what one believes to be true (2002, 96). While originally spelled out in virtue-theoretic terms, we can easily transport this distinction to the properties of asserting true beliefs and saying only what one believes to be true. On this account, honesty is a necessary but not sufficient condition for truthful language use. (This distinction is analogous to the distinction between truthfulness as the property of an utterance being true (Evans et al. 2021), and honesty as only saying what one believes is true [Friedl, Ward et al., manuscript]).

Yet, merely avoiding a lie is not sufficient to make an utterance honest. This intuition is again appreciated by Williams (2002, 97), according to which trustworthiness requires more than just the avoidance of lying. And more recent work, that focuses on the speaker's motivations in explicating the cooperative conditions of successful communication, has drawn the attention to how even in the absence of a lie, an utterance can be deceptive, more specifically misleading (Saul 2013; Stokke 2013, 2018; Viebahn 2021). This work has largely focussed on defining a misleading utterance with explicit reference to speaker intentions (Carson 2012; Saul 2013; Stokke 2013, 2018). Some argue that insincere language necessitates the intention to deceive (Carson 2012; Saul 2013), while others argue that it is the absence of a certain intention that is necessitated by insincere language use (Stokke 2013, 2018).

One desideratum for a notion of honesty adequate for addressing ELK is that it needs to accommodate that a speaker's utterance may be dishonest, even if she does not have the intention to deceive (e.g., representing the best guess for the value of a variable and still being dishonest). Here, accounts like Stokke (2013), who defines insincerity of declarative sentences as follows, seem more adequate as a starting point:

A declarative utterance u by a speaker A is insincere with respect to a QUD q if and only if, in making u , A does not consciously intend to:

(I1) communicate an answer p to q such that A mentally assents to p , and

(I2) not communicate some proposition p' such that A does not assent to p' . [129]

This idea builds on the Stalnakerian framework of a conversation being a purpose-oriented cooperation to inquire about the Big Question of the conversation. This account of the question-oriented cooperative language game has been formalised by Roberts (2012) as the Question Under Discussion (QUD), where the interlocutors are cooperative in asserting a sentence iff they aim to answer the QUD with their assertion (e.g., 'Will it rain this evening?'). An utterance's sincerity is then determined by whether it

is performed with the intention to answer the QUD. On this account, if the speaker does not utter p with the intention to make the listener believe p (which she has assented to), her utterance is insincere ('assented' can be operationalised as 'accepted' where the agent acts as if she believes the proposition specifically in the context of the conversation, see e.g., Dinges (2024)).

Stokke's account focuses on defining insincerity (dishonesty) as opposed to sincerity (honesty). Yet, the intuitive account of an honest utterance we have in mind is structurally analogous to this approach, so it is easy to see how it can be extended into a positive account of an honest utterance, where the speaker must have a certain intention with which she utters a sentence: An intention to bring about that the hearer forms a true belief with regard to the question she has asked.

There have also been approaches to defining honesty more explicitly. Wilson (2018), for instance, argues for a motivational account of honesty, according to which honesty necessitates the motivation to avoid deception. Similarly, Sbardolini (2022, 981) explicitly defines honesty as "knowingly saying something one believes to be true relative to the world of interest, which may or may not be actual" and dishonesty as "knowingly saying something one believes to be false relative to the world of interest." Similarly, Stalnaker's (2006, 92–93) notion of credibility also touches on what an honest utterance would look like:

A message m is credible if, and only if, a rational speaker prefers that the listener believes m just in case m is true (relative to the world of interest in the conversation). (Sbardolini 2022, 992)

While the philosophical literature has fruitfully moved towards defining sincere and insincere language use with explicit reference to the reason for which a speaker utters a sentence (i.e. motivational reason/intention in action) and whether the reason is a cooperative one, the exact workings of such speaker intentions/reasons need more careful attention, specifically with regards to what is the content of the speaker intention: e.g., avoiding deception, making the listener form a true belief about the question they have asked, making the listener believe what the speaker herself believes, etc.

Carson, Thomas L. 2012. *Lying and Deception: Theory and Practice*. Oxford University Press.

Dinges, Alexander. 2024. "Assertion and Certainty." *The Philosophical Quarterly* 74 (1): 169–86. <https://doi.org/10.1093/pq/pqad022>.

Evans, Owain, Owen Cotton-Barratt, Lukas Finnveden, Adam Bales, Avital Balwit, Peter Wills, Luca Righetti, and William Saunders. 2021. "Truthful AI: Developing and Governing AI That Does Not Lie." arXiv. <https://doi.org/10.48550/arXiv.2110.06674>.

Grice, H. P. 1975. "Logic and Conversation." In *The Semantics-Pragmatics Boundary in Philosophy*, edited by M. Ezcurdia and R. J. Stainton, 47–59. Peterborough: Broadview Press.

Roberts, Craig. 2012. "Information Structure in Discourse: Towards an Integrated Formal Theory of Pragmatics." *Semantics & Pragmatics* 5 (6): 1–69. <https://doi.org/10.3765/sp.5.6>.

- Saul, Jennifer Mather. 2013. *Lying, Misleading, and What Is Said: An Exploration in Philosophy of Language and in Ethics*. Oxford University Press.
- Sbardolini, Giorgio. 2022. “Is Honesty Rational?” *The Philosophical Quarterly* 72 (4): 979–1001. <https://doi.org/10.1093/pq/pqab071>.
- Stalnaker, Robert. 2006. “Saying and Meaning, Cheap Talk and Credibility.” In *Game Theory and Pragmatics*, edited by A. Benz, G. Jäger, and R. van Rooij, 83–100. London: Palgrave Macmillan.
- Stokke, Andreas. 2013. “Lying and Asserting.” *Journal of Philosophy* 110 (1): 33–60. <https://doi.org/10.5840/jphil2013110144>.
- . 2018. *Lying and Insincerity*. Oxford University Press. <https://doi.org/10.1093/oso/9780198825968.001.0001>.
- Viebahn, Emanuel. 2021. “The Lying-Misleading Distinction: A Commitment-Based Approach.” *Journal of Philosophy* 118 (6): 298–319. <https://doi.org/10.5840/jphil2021118621>.
- Williams, Bernard. 2002. *Truth and Truthfulness: An Essay in Genealogy*. Princeton University Press.
- Wilson, A. T. 2018. “Honesty as a Virtue.” In *Connecting Virtues*, edited by M. Croce and M. S. Vaccarezza. <https://doi.org/10.1002/9781119525660.ch4>.